

Classification of Paintings

Project: Deep Learning TensorFlow (2021.01, opencampus)
Nils Berns, John Kimani

Objectives

Goal

Implementation of a deep learning convolution neural network to classify painting of various influential painters.

We do this by achieving the following objectives:

1. Aggregate data from different sources to achieve objectives (2 – 4)
2. Match the painting by artist.
3. Differentiate between real and fraudulent paintings.
4. Differentiate between painting-stylised photos with actual painting.



Data

Data sets

„Best Artworks of All Time“ [1]

<https://www.kaggle.com/ikarus777/best-artworks-of-all-time>

Collection of paintings of the 50 most influential artists of all time.
8446 images.



„<https://thisartworkdoesnotexist.com/>“ [2]

With every request (1 per sec) an artificial painting is generated
Style is quite modern/contemporary.
Downloaded 1000 images.

„bored humans AI paintings“ [3]



<https://boredhumans.b-cdn.net/art/>

Database of artificial paintings generated with styleGAN2
Downloaded 1000 images of ~5000 with.

„Paired landscape and Monet-Stylized image“ [4]

<https://www.kaggle.com/shcsteven/paired-landscape-and-monetstylised-image>

Used only the stylised images.
1030 images.



Challenges

with “Best Artworks of All Time”

- Most artists have less than 100 paintings others more than 500.
- Artists from the same genre share more features.
- Artist's style changes over time.
- Trade-off between computational costs and the image size required for capturing all style features.

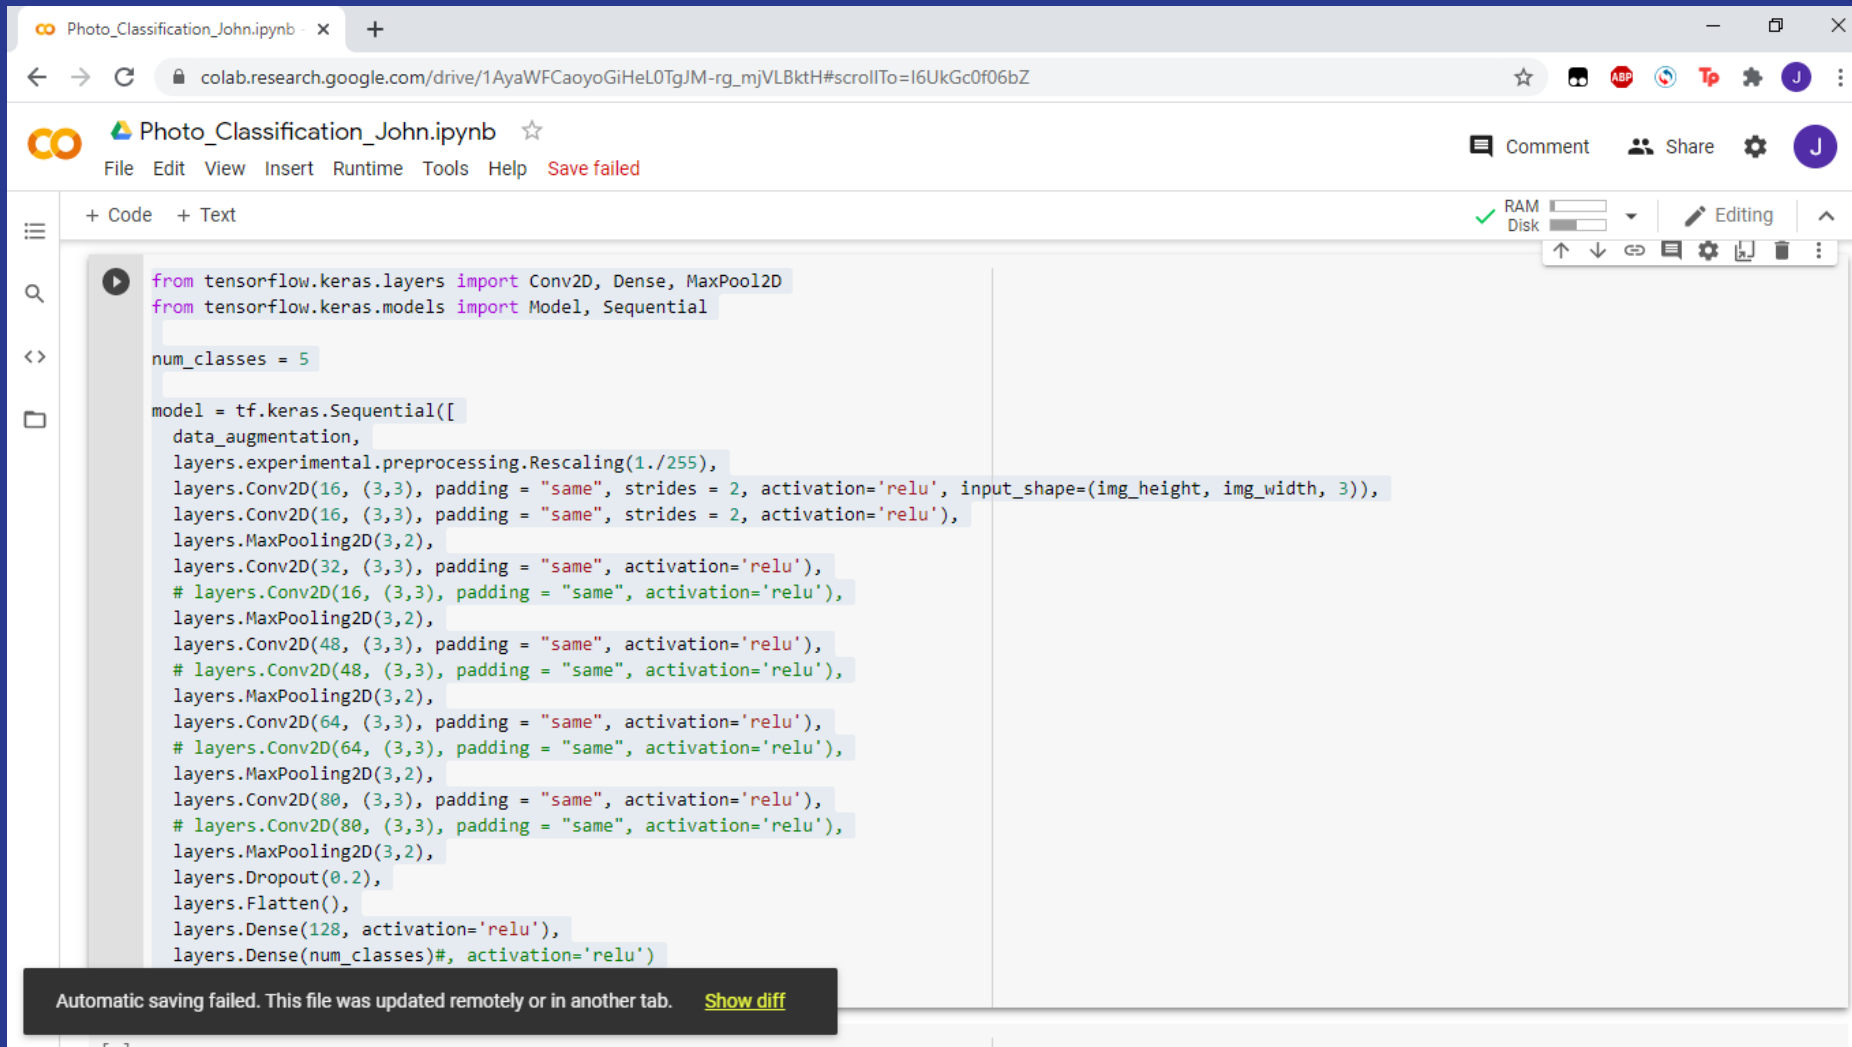
Implementation

- Choices:
- Minimum number of paintings an artist has to have.
- Inclusion of an additional class with fake/artificial paintings (max. resolution 512x512).
- Downloading and organising the data takes a while during first run.
- 80% are assigned to the training set, 10% to development and test set each.



Technical Environment

Implementation



```
from tensorflow.keras.layers import Conv2D, Dense, MaxPool2D
from tensorflow.keras.models import Model, Sequential

num_classes = 5

model = tf.keras.Sequential([
    data_augmentation,
    layers.experimental.preprocessing.Rescaling(1./255),
    layers.Conv2D(16, (3,3), padding = "same", strides = 2, activation='relu', input_shape=(img_height, img_width, 3)),
    layers.Conv2D(16, (3,3), padding = "same", strides = 2, activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Conv2D(32, (3,3), padding = "same", activation='relu'),
    # layers.Conv2D(16, (3,3), padding = "same", activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Conv2D(48, (3,3), padding = "same", activation='relu'),
    # layers.Conv2D(48, (3,3), padding = "same", activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Conv2D(64, (3,3), padding = "same", activation='relu'),
    # layers.Conv2D(64, (3,3), padding = "same", activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Conv2D(80, (3,3), padding = "same", activation='relu'),
    # layers.Conv2D(80, (3,3), padding = "same", activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)#, activation='relu')
])
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

Colab for coding

- TensorFlow 2.4.0
- Shairing colab via google Drive in Colab



Model

Data Augmentation

```
data_augmentation = tf.keras.Sequential(  
    [  
        layers.experimental.preprocessing.RandomFlip("horizontal", input_shape=(img_height, img_width, 3)),  
        layers.experimental.preprocessing.RandomRotation(0.1),  
        layers.experimental.preprocessing.RandomZoom(0.1),  
    ]  
)
```

Model Architecture

```
from tensorflow.keras.layers import Conv2D, Dense, MaxPool2D
from tensorflow.keras.models import Model, Sequential
num_classes = 5
model = tf.keras.Sequential([
    data_augmentation,
    layers.experimental.preprocessing.Rescaling(1./255),
    layers.Conv2D(16, (3,3), padding = "same", strides = 2, activation='relu', input_shape=(img_height, img_width, 3)),
    layers.Conv2D(16, (3,3), padding = "same", strides = 2, activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Conv2D(32, (3,3), padding = "same", activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Conv2D(48, (3,3), padding = "same", activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Conv2D(64, (3,3), padding = "same", activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Conv2D(80, (3,3), padding = "same", activation='relu'),
    layers.MaxPooling2D(3,2),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes
])
```

Optimiser

```
model.compile(  
    optimizer='adam',  
    loss=tf.losses.SparseCategoricalCrossentropy(from_logits  
=True),  
    metrics=['accuracy'])
```

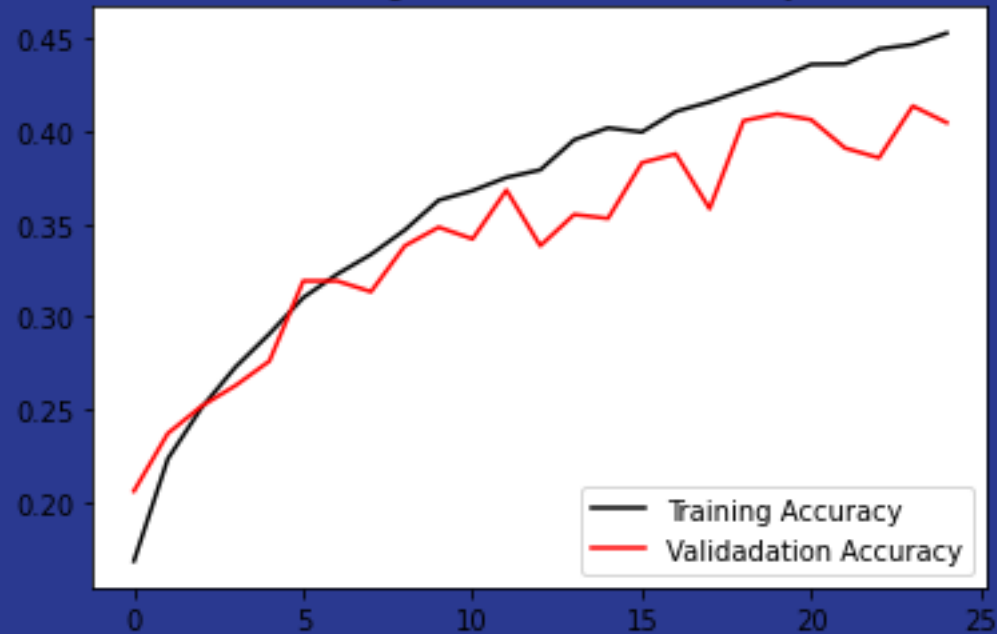


Performance and Result

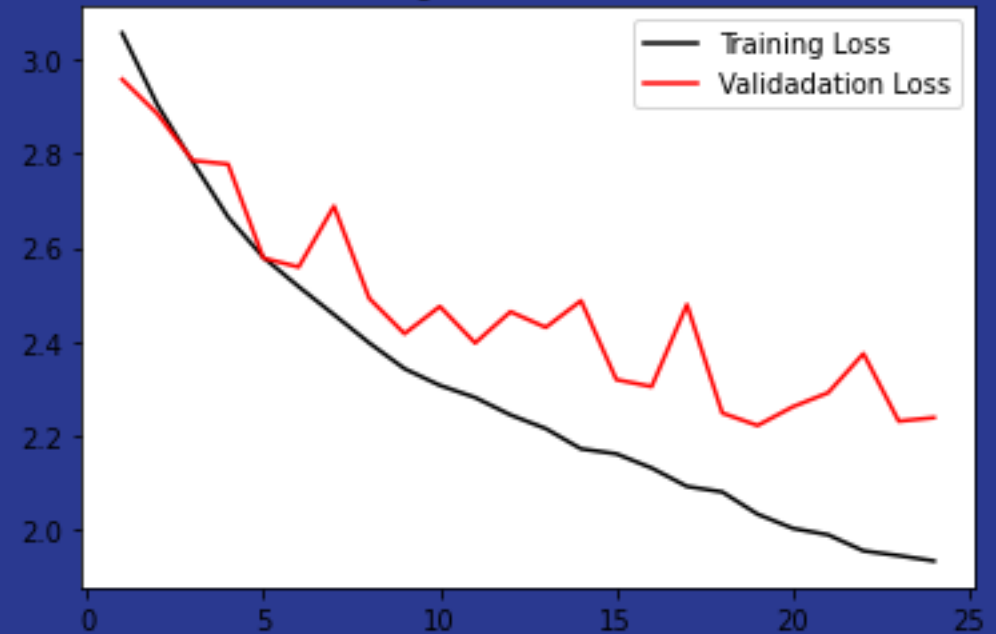
50 artists and Monet Stylised image

Model with data augmentation and 1 drop-out layer

Training and validation accuracy





Training and validation loss





4 artists and Monet Stylised image

Model with data augmentation and drop-out layer





4 most populated „artists“ and Monet Stylised image

Model with no data augmentation and drop out layer



4 most populated „artists“ and Monet Stylised image

Model with data augmentation and no Drop out layer



Comparison

Number of artists	Max. training accuracy	Max. validation accuracy
50 + Monet stylised images	~ 50%	~ 50%
11	~ 75%	~ 70%
4+ Monet stylised images	~ 95%	~ 80%

Best setting

- 5 most populated „artists“
- Augmented Data
- No drop-out layer

Conclusions

Conclusions

- Most artists in the data-set have too few paintings
 - fewer artists with more paintings increase the accuracy
 - data augmentation increased the performance
- The model architecture does not have to be complex
- Overfitting can occur
- Adding the fake class increases the accuracy by ~5%
- Size of the data-set becomes an issue in Colab, when the resolution is too high (even with 5 classes).
- Random search for best hyperparameters was terminated by Colab due to the duration (10 sets, not in parallel, image size 128x128).

Outlook

- Evaluate more performance metrics
- Plot the outputs of the convolutional layers
- Use training and test set only
- Use less popular painters with more paintings
- Perhaps avoid painters with too similar style
- Prefer painters with a consistent style
- Search for a better peer model architecture
- Use transfer learning with a pre-trained model
- Parallelise the hyperparameter search
- Increase the image resolution to capture more features